

Processor controls hardware and executes instructions.
Memory stores data and code.
I/O components read and write from I/O devices.
System bus interconnects different components and provides communication between them.
OS must make **efficient** use of available resources and must **share resources** (CPU cores, caches, RAM, GPUs, clocks, timers, persistent storage) among multiple users.
Hardware abstraction layer — converts raw hardware into a usable computer system by keeping details hidden from user-level software.
Must support several simultaneous parallel activities, and switch activities at arbitrary times, so must ensure **safe concurrency**.

MULTIPROCESSORS (Windows, macOS, Linux) have many CPU cores and CPUs.

SERVER OS (Solaris, FreeBSD, Linux) shares hardware/
MAINFRAMES have bespoke hardware, limited workload.

SMARTPHONES have power-efficient CPUs.
EMBEDDED OS (QNX, VxWorks) are used for home utilities, and only run trusted software.

REAL-TIME OSes are time-oriented (not performance/IO).

SENSOR NETWORK OS (TinyOS) is resource efficient.
The OS **kernel** is a privileged program that implements OS functionality.

Monolithic kernels are a single black box with all functionality. Only executable with its own address space.
Microkernels have as little functionality in the kernel as possible, functionality passed to a user-level server. Kernel does IPC between servers — separate ones for I/O, file access, memory management, etc. less overhead. Less error-prone and servers have clean interfaces.

Hybrid kernels combine features of both.
LINUX is a variant of Unix. **Interrupt handlers** are the primary means to interact with devices. I/O scheduler orders disk operations. Supports static in-kernel components and dynamically loadable modules.

NTOS practices **systematic calls** in **WINDOWS**. Programs run on top of **Windows API**. **Windows API** — NTOS loaded from **ntoskrnl.exe** at boot, consists of **executive** and **kernel** layers. **Device drivers** loaded into kernel memory. NT structured like a **microkernel**, most components run in same address space. **Hardware Abstraction Layer** abstracts out BIOS config, CPU types, etc...

PROCESSES allow a single processor to run multiple programs **"simultaneously"**. Provides concurrency, isolation between programs, allows better resource utilization, simplifies programming.
Pseudo concurrency single physical processor is switched between processes by interleaving.
Real concurrency utilises multiple physical processors.

On a **context switch**, the processor switches from executing process A to executing process B. May occur in response to events/interrupts, or as a result of periodic scheduling. Process can **switch later** to C, so no information needed about what is stored in **process descriptor** / **control block**, kept in the **process table**.

Process Control Block — process has its own VM, should store PC, page table register, stack pointer, ..., process management info, file management info.
Context switches are expensive, so avoid unnecessary context switching. **Indirect context switching** process A context — perturbation of memory caches.
flushing TLB & pipeline flushing

UNIX allows processes to create hierarchies. **Windows** has no notion of this.
int fork(void) Creates new child process, to be joined by parent. By default it's a copy-on-write of parent's resources of parent's process image. Returns child PID in parent process and 0 in child process. No child created, and I returned to parent on error.

execve(char *path, char *argv[], char *envp[]) — Changes process image and runs new process. Lots of useful wrappers.
int waitpid(int pid, int status, int options) — Suspends execution of calling process until process pid terminates normally or signal received. Can wait for more than one child.

Simple design philosophy, so process API made up of basic blocks that can be easily combined.

CreateProcess() (equivalent of **fork()** & **exec()**) has 10 parameters!

Termination — normally process completes execution of body, **call(MT)**. After exit — runs into error and is killed except if **aborted**, another process has overruled execution. Some in an endless loop.

void exit(int status) — Terminates a process. Never returns in the calling process.
int kill(int pid, int sig) — sends signal sig to process pid.

SIGNALS are an IPC mechanism
SIGABRT: Abort signal from assert
SIGFPE: Floating point exception
SIGKILL: Kill signal
SIGPIPE: Broken pipe; write to pipe with no readers
SIGALRM: Timer signal from alarm
SIGTERM: Termination signal
Generated when an exception occurs, if it has permission.

Generated when an exception occurs, when certain key combinations typed in terminal, or programmatically using **mt kill**.

Default action for signals is to terminate process, but receiving process may ignore/handle it instead (unless **SIGKILL** or **SIGSTOP**).
Pipes method of connecting stdout of one process to stdin of another. Read or unimpl. Sender should close the read end, receiver should close the write end.
Named Pipes / **FIFOs** are persistent pipes that survive processes which created them. Stored on file system. Any process can open it like a regular file.
int pipe(int fd[2]) — returns two file descriptors: read end(fd[0]), write end(fd[1]).

THREADS are execution streams that share the same address space. **Multithreading** — each process can contain multiple threads.

Useful for applications which contain multiple parallel activities that share the same data. Processes are too heavyweight. It is difficult to coordinate between address spaces, and expensive to context switch between them & create/destroy activities.

Shared address space can lead to memory corruption and concurrency bugs. How to handle **locking and signals?**

PTHREADS implemented by most Unix systems.

int pthread_create(...) — Creates a new thread.

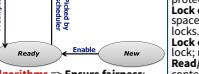
void pthread_exit(void *ptr) — Terminates the thread and makes value **ptr** available to any successful join with terminating thread. Called implicitly when the

thread's start routine returns.
int pthread_yield(void) — Releases CPU to let another thread run. Returns 0 on success (always), or error code.
int pthread_join(...) — Blocks until thread terminates.
User-level threads — OS kernel not aware of threads, each thread manages its own threads. Better performance, and each application can have its own scheduling algorithm. But blocking system calls (e.g. during page fault) stop **all threads** in process.

Hybrid approaches use kernel threads and multiplex user-level threads onto some of these.

SCHEDULING

If multiple processes are ready, which one should be run?



Goals of Scheduling Algorithms — Ensure **fairness**: Comparable processes should get comparable services; Avoid **idle postponement**: No process should starve; Enforce **policy**: E.g. priorities; Maximize **resource utilization**: CPU, I/O devices; Minimize **overhead**: From context switches, lock contention decisions.

For Batch Systems => **Throughput**: Jobs per unit of time; **Turnaround time**: Time between job submission & completion.

For Real-time Systems => **Response time**: Time between request issued and first response.

Non-preemptive scheduling lets processes run until it blocks / releases CPU. **Preemptive** scheduling lets processes run for a maximum amount of time.

FIRST-COME FIRST-SERVE — No indefinite postponement. Easy implementation. But suffers from "head-of-line blocking" if long job followed by many short jobs.

ROUND-ROBIN: Process runs until it blocks or time quantum is reached, and goes to ready queue. Poor turnaround time when jobs have similar run-times.
Good avg. turnaround time when different runtimes.
RR quantum (time slice) => should be larger than context switch cost, but still provide decent response times.

HYBRID SCHEDULING — Combines FCB and RR. Context switches are expensive, so avoid unnecessary context switching. Within process, use RR. Between processes, use FCB.

PRIORITY SCHEDULING — Jobs run based on their priority. Priority can be static or dynamic.

Generally, favour short and I/O-bound jobs. Quickly determine the nature of job and adapt to changes.

MULTILEVEL FEEDBACK QUEUES — Form of priority scheduling. One queue for each priority level. Within each queue, can use round robin or RR. Need to determine current nature of jobs. To prevent starvation of lower-priority jobs. Feedback comes from periodically recomputing priorities: e.g. based on how much CPU they used recently, increase job's priority. Or based on average weighted moving average. Not very flexible, does not react quickly to changes, cheating is a concern.

LOTTERY SCHEDULING — Jobs receive lottery tickets for various resources — random ticket chosen at each scheduling decision. Number of lottery tickets is meaningful as a weight. Highly responsive with no starvation. Jobs can exchange tickets. But unpredictable response time.

How do processes **SYNCHRONISE** their operations to perform a task?

Critical section — processes access a shared resource. **Synchronisation mechanism** required at entry/exit of critical section.

Race condition occurs when multiple threads read and write shared data within critical section => their execution interleaved non-deterministically.

Assume **sequential consistency** throughout — operations of each thread appear in program order, operations of all threads executed in sequential order atomically.

Synchronisation mechanism required at entry/exit of critical section => mediate concurrent access & prevent race conditions.

An **atomic operation** is a sequence of 1+ statements that are appear to be indivisible.

TEST & SET LOCK is an atomic instruction **TS(L)** to 1 and returns old value. Locks using busy waiting (e.g. **if (TS(L) != 0) /> wait();**) are **spin locks**.

DISALLOWED INTERRUPTS only works on single-processor systems. Code-behind/buggy processes may never release the CPU.

STRICT ALTERNATION fails to meet rules of mutual exclusion = if T_0 's non-crit. section is **long enough**, it will prevent T_1 from entering crit. section. $\infty \infty$

T₀

T₁

SIGKILL or **SIGSTOP**

Default action for signals is to terminate process, but receiving process may ignore/handle it instead (unless **SIGKILL** or **SIGSTOP**).

PIPES method of connecting stdout of one process to stdin of another. Sender should close the read end, receiver should close the write end.

Named Pipes / **FIFOs** are persistent pipes that survive processes which created them. Stored on file system. Any process can open it like a regular file.

int pipe(int fd[2]) — returns two file descriptors: read end(fd[0]), write end(fd[1]).

PTHREADS implemented by most Unix systems.

int pthread_create(...) — Creates a new thread.

void pthread_exit(void *ptr) — Terminates the thread and makes value **ptr** available to any successful join with terminating thread. Called implicitly when the

thread's start routine returns.
Mutual exclusion ensures if a process executes its critical section, no other process can be executing it. 1. Many processes may be simultaneously inside a critical section.
2. No process running outside critical section may prevent other processes entering the critical section.
3. No process requiring access to its critical section can be delayed forever.
4. All processes made about the relative speed of processes.
A commonly desired feature of synchronisation mechanisms.

Example implementations => disabling interrupts, semaphores w/ inc/dec counter-1.

LOCK/MUTEX => mutual exclusion mechanism owned by actively excluding thread; lock is usually user/process object while mutex is same thing but OS/Kernel object.

lock granularity => the amount of data a lock is protecting.

Lock overhead => cost of using locks, e.g. memory space, initialisation, time required to acquire/release lock.

Lock contention => number of processes waiting for lock; more contention => less parallelism.

Read/Write Locks can be used to minimize lock contention & maximize concurrency => exclusive access mode but multiple threads can acquire in read mode.

SEMAPHORES — processes cooperate by means of signals. Special variables, accessible via (atomic) **down(s)**, **up(s)**, **init(s, i)**. Two private components — a **counter** and a **queue** of processes waiting on the semaphore.

In a **binary semaphore**, counter initialised to -1.

```
int(s): ai := counter(s) + 1; down(s): : If counter(s) <= 0 then
    if (ai == 0) then
        if (queue(s) == empty) then
            if (counter(s) > 0) then
                add(s) to queue(s)
            else
                remove one process in queue(s); suspend current process
                else
                    counter(s) = counter(s) + 1
```

MONITORS on COND-VARS:

Monitors regulate access to shared data & consist of (implicit) **monitor lock** and 1+ **cond-vars**; **entry** procedures called from outside the monitor, **internal** procedures only callable from **monitor procedures**, processes can only call entry procedures — only one process can be in the monitor at one time

Cond-vars associated with high-level conditions; if **cond-var** signalled w/ no one waiting for it then the process is **locked**; has operations:

wait(c) => release monitor lock & wait for c to be signalled.

signal(c) => wake up one process waiting for c

broadcast(c) => wake up all processes waiting for c

Monitors are a language construct and not supported by C (supported by PintOS via explicit monitor locks). Java supports variant of monitors via **synchronized** methods, but they lack **condition** variables, i.e. no **wait()**/**notify()**

Producer/consumer example w/ monitors ¶ ¶

Monitor **ProduceConsumer**

```
condition not full, not_empty;
integer count = 0;
```

```
entry procedure insert(item)
    while (count == N) wait(not_full);
    insert_item(item); count++;
    signal(not_empty);
```

```
entry procedure remove(item)
    while (count == 0) wait(not_empty);
    remove_item(item); count--;
    signal(not_full);
```

end monitor

^{^ producer inserts & consumer removes items w/ max. capacity N => only insert/remove if mutual exclusion achieved => only insert if non-full & remove if non-empty.}

How do processes **SYNCHRONISE** their operations to perform a task?

Critical section — processes access a shared resource. **Synchronisation mechanism** required at entry/exit of critical section.

Race condition occurs when multiple threads read and write shared data within critical section => their execution interleaved non-deterministically.

Assume **sequential consistency** throughout — operations of each thread appear in program order, operations of all threads executed in sequential order atomically.

Synchronisation mechanism required at entry/exit of critical section => mediate concurrent access & prevent race conditions.

An **atomic operation** is a sequence of 1+ statements that are appear to be indivisible.

TEST & SET LOCK is an atomic instruction **TS(L)** to 1 and returns old value. Locks using busy waiting (e.g. **if (TS(L) != 0) /> wait();**) are **spin locks**.

DISALLOWED INTERRUPTS only works on single-processor systems. Code-behind/buggy processes may never release the CPU.

STRICT ALTERNATION fails to meet rules of mutual exclusion = if T_0 's non-crit. section is **long enough**, it will prevent T_1 from entering crit. section. $\infty \infty$

T₀

T₁

SIGKILL or **SIGSTOP**

Default action for signals is to terminate process, but receiving process may ignore/handle it instead (unless **SIGKILL** or **SIGSTOP**).

PIPES method of connecting stdout of one process to stdin of another. Sender should close the read end, receiver should close the write end.

Named Pipes / **FIFOs** are persistent pipes that survive processes which created them. Stored on file system. Any process can open it like a regular file.

int pipe(int fd[2]) — returns two file descriptors: read end(fd[0]), write end(fd[1]).

PTHREADS implemented by most Unix systems.

int pthread_create(...) — Creates a new thread.

void pthread_exit(void *ptr) — Terminates the thread and makes value **ptr** available to any successful join with terminating thread. Called implicitly when the

thread's start routine returns.
Busy waiting => synch. mechanism for which the waiting thread remains in running state (e.g. spin locks) which wastes CPU cycles, and runs into **PRIORITY INVERSION PROBLEM** if higher-priority H is busy-waiting on lock held by lower-priority L, then H is always scheduled meaning it cannot make progress, i.e.

alternatively, thread T_B calls enter/leave_critical(); ¶ ¶

```
int interested[2] = {0, 0}; // thread is 0 or 1
void enter_critical(int thread)
    int other = 1 - thread;
    if (interested[thread] == 0) {
        interested[thread] = 1;
        if (interested[other] == 0) {
            interested[other] = 1;
            turn = other;
            while (turn == other && !interested[other])
                loop();
        }
    }
}
```

PETERSON'S SOLUTION fixes issues w/ strict alternation; thread T_B calls enter/leave_critical(); ¶ ¶

```
int interested[2] = {0, 0}; // thread is 0 or 1
void leave_critical(int thread)
    int other = 1 - thread;
    if (interested[thread] == 1) {
        interested[thread] = 0;
        if (interested[other] == 1) {
            interested[other] = 0;
            turn = other;
            while (turn == other && !interested[other])
                loop();
        }
    }
}
```

alternatively, thread T_B calls enter/leave_critical(); ¶ ¶

```
int interested[2] = {0, 0}; // thread is 0 or 1
void leave_critical(int thread)
    int other = 1 - thread;
    if (interested[thread] == 1) {
        interested[thread] = 0;
        if (interested[other] == 1) {
            interested[other] = 0;
            turn = other;
            while (turn == other && !interested[other])
                loop();
        }
    }
}
```

alternatively, thread T_B calls enter/leave_critical(); ¶ ¶

```
int interested[2] = {0, 0}; // thread is 0 or 1
void leave_critical(int thread)
    int other = 1 - thread;
    if (interested[thread] == 1) {
        interested[thread] = 0;
        if (interested[other] == 1) {
            interested[other] = 0;
            turn = other;
            while (turn == other && !interested[other])
                loop();
        }
    }
}
```

alternatively, thread T_B calls enter/leave_critical(); ¶ ¶

```
int interested[2] = {0, 0}; // thread is 0 or 1
void leave_critical(int thread)
    int other = 1 - thread;
    if (interested[thread] == 1) {
        interested[thread] = 0;
        if (interested[other] == 1) {
            interested[other] = 0;
            turn = other;
            while (turn == other && !interested[other])
                loop();
        }
    }
}
```

alternatively, thread T_B calls enter/leave_critical(); ¶ ¶

```
int interested[2] = {0, 0}; // thread is 0 or 1
void leave_critical(int thread)
    int other = 1 - thread;
    if (interested[thread] == 1) {
        interested[thread] = 0;
        if (interested[other] == 1) {
            interested[other] = 0;
            turn = other;
            while (turn == other && !interested[other])
                loop();
        }
    }
}
```

alternatively, thread T_B calls enter/leave_critical(); ¶ ¶

```
int interested[2] = {0, 0}; // thread is 0 or 1
void leave_critical(int thread)
    int other = 1 - thread;
    if (interested[thread] == 1) {
        interested[thread] = 0;
        if (interested[other] == 1) {
            interested[other] = 0;
            turn = other;
            while (turn == other && !interested[other])
                loop();
        }
    }
}
```

alternatively, thread T_B calls enter/leave_critical(); ¶ ¶

```
int interested[2] = {0, 0}; // thread is 0 or 1
void leave_critical(int thread)
    int other = 1 - thread;
    if (interested[thread] == 1) {
        interested[thread] = 0;
        if (interested[other] == 1) {
            interested[other] = 0;
            turn = other;
            while (turn == other && !interested[other])
                loop();
        }
    }
}
```

alternatively, thread T_B calls enter/leave_critical(); ¶ ¶

```
int interested[2] = {0, 0}; // thread is 0 or 1
void leave_critical(int thread)
    int other = 1 - thread;
    if (interested[thread] == 1) {
        interested[thread] = 0;
        if (interested[other] == 1) {
            interested[other] = 0;
            turn = other;
            while (turn == other && !interested[other])
                loop();
        }
    }
}
```

alternatively, thread T_B calls enter/leave_critical(); ¶ ¶

```
int interested[2] = {0, 0}; // thread is 0 or 1
void leave_critical(int thread)
    int other = 1 - thread;
    if (interested[thread] == 1) {
        interested[thread] = 0;
        if (interested[other] == 1) {
            interested[other] = 0;
            turn = other;
            while (turn == other && !interested[other])
                loop();
        }
    }
}
```

alternatively, thread T_B calls enter/leave_critical(); ¶ ¶

```
int interested[2] = {0, 0}; // thread is 0 or 1
void leave_critical(int thread)
    int other = 1 - thread;
    if (interested[thread] == 1) {
        interested[thread] = 0;
        if (interested[other] == 1) {
            interested[other] = 0;
            turn = other;
            while (turn == other && !interested[other])
                loop();
        }
    }
}
```

alternatively, thread T_B calls enter/leave_critical(); ¶ ¶

```
int interested[2] = {0, 0}; // thread is 0 or 1
void leave_critical(int thread)
    int other = 1 - thread;
    if (interested[thread] == 1) {
        interested[thread] = 0;
        if (interested[other] == 1) {
            interested[other] = 0;
            turn = other;
            while (turn == other && !interested[other])
                loop();
        }
    }
}
```

alternatively, thread T_B calls enter/leave_critical(); ¶ ¶

```
int interested[2] = {0, 0}; // thread is 0 or 1
void leave_critical(int thread)
    int other = 1 - thread;
    if (interested[thread] == 1) {
        interested[thread] = 0;
        if (interested[other] == 1) {
            interested[other] = 0;
            turn = other;
            while (turn == other && !interested[other])
                loop();
        }
    }
}
```

alternatively, thread T_B calls enter/leave_critical(); ¶ ¶

```
int interested[2] = {0, 0}; // thread is 0 or 1
void leave_critical(int thread)
    int other = 1 - thread;
    if (interested[thread] == 1) {
        interested[thread] = 0;
        if (interested[other] == 1) {
            interested[other] = 0;
            turn = other;
            while (turn == other && !interested[other])
                loop();
        }
    }
}
```

alternatively, thread T_B calls enter/leave_critical(); ¶ ¶

```
int interested[2] = {0, 0}; // thread is 0 or 1
void leave_critical(int thread)
    int other = 1 - thread;
    if (interested[thread] == 1) {
        interested[thread] = 0;
        if (interested[other] == 1) {
            interested[other] = 0;
            turn = other;
            while (turn == other && !interested[other])
                loop();
        }
    }
}
```

alternatively, thread T_B calls enter/leave_critical(); ¶ ¶

```
int interested[2] = {0, 0}; // thread is 0 or 1
void leave_critical(int thread)
    int other = 1 - thread;
    if (interested[thread] == 1) {
        interested[thread] = 0;
        if (interested[other] == 1) {
            interested[other] = 
```