



<p><b>Level 5 (Block-level Distributed XOR)</b> — Distributed parity information. Some potential for write concurrency, lower efficiency/redundancy trade-off.</p>	<p><b>DISK CACHE</b> — Main memory buffer contains copy of disk sectors. Uses finite space, so need a <b>replacement policy</b> when buffer full.</p>
<p><b>LRU</b> — Cache consists of stack of blocks, remove block cache longest with no references. Use a stack of <b>pointers</b> to track blocks that has experienced fewest references. Counter associated with each block, incremented each time block accessed.</p>	<p><b>Frequency-Based</b> — Divide LRU stack into new/old sections. Block referenced → move to top of stack. Increment reference count if not already in new. Replace old block with lowest frequency.</p>
<p><b>LRU</b> — Keeps block that has experienced fewest references. Counter associated with each block, incremented each time block accessed.</p>	<p>To prevent blocks “aging out” too quickly, could use three sections and still only replace blocks from <i>old</i>.</p>
<p><b>FILE SYSTEMS</b> — we want:</p> <ul style="list-style-type: none"><li>Long term, nonvolatile, online storage.</li><li>Sharing of data.</li><li>Organisation and management of data.</li></ul>	<p><b>FILE</b> — named collection of data of arbitrary size. Named with an extension based off what type of file they are.</p> <p><b>File Functions</b> — <i>Create, Delete, Open, Close, Read Write, Reposition/Seek, Truncate, Rename, Read attributes, Write attributes</i>.</p> <p><b>File System Management Functions</b> — Logical name to physical disk address translation. Management of disk space. File locking for exclusive access, performance optimisation, Protection against system failure, Security.</p> <p><b>File attributes</b> may be held within given directory, system <b>base attributes</b> — file metadata, <b>extended attributes</b> — additional information (volume/start address/size/used/size allocated), <b>access control information</b> (owner/authentication/permited actions), <b>usage information</b> (creation timestamp/modification/last read/last archived/expiry date).</p> <p>File attributes can be accessed with the <b>stat syscall</b>, Returns information about specified file in <b>struct stat</b>.</p> <p><b>ORGANISATION</b></p> <p><b>Dynamic space management</b> — space allocated in blocks as file size natural variable. Large block size wastes space for small files; small block sizes waste space for large files. Various methods for allocating blocks.</p> <p><b>Contiguous File Allocation</b> — Place file data at contiguous addresses on storage device. Successive logical records physically adjacent. Can lead to external fragmentation, poor performance if files grow/shrink over time.</p> <p><b>Block Linkage (Chaining)</b> — Chain must be searched sequentially for data. Wastes space as each block has space in each block. Insertion/deletion by modifying pointer in previous block. Large block sizes result in internal fragmentation. Small block sizes — data spread across multiple blocks, poor performance due to many seeks.</p> <p><b>Block Allocation Table</b> — stores pointers to file blocks. <b>Free block bitmap</b> — Table stores one bit, cached in memory for performance. Reduces number of lengthy seeks to access given record (but files become fragmented).</p> <p>Each file has 1+ <b>index block</b> — contains list of pointers that point to data blocks. May change by request, last pointer to store pointers to data blocks. Searching may take place in the blocks themselves; if they are near corresponding data blocks → quick access to data.</p> <p><b>Inodes</b> are <b>index blocks</b>. On file open, OS opens <b>inode table</b>, structured as an inode on disk but including device specific info. Inodes are shared between processes with opened file, major/minor device number.</p> <p>Use a <b>free list</b> — linked list of locations of free blocks → to manage storage device's free space. Low overhead to perform maintenance operations, file likely to be allocated from non-contiguous blocks.</p> <p><b>Bitmap</b> contains one bit in memory for each disk block, indicating whether in use. Can quickly determine available <b>contiguous</b> blocks at certain locations, but, unlike free list, may need to search entire bitmap to find free block.</p> <p><b>Boot block</b> — LAYOUT — Fixed disk layout with inodes — <b>superblock</b> contains crucial info about FS.</p> <p><b>Super block</b> — <b>Directory</b> maps symbolic file names to physical disk location.</p> <p><b>Inode bitmap</b> — In hierarchical file system, root indicates where on disk root directory begins, points to various directories. File names only need to be unique within given directory.</p> <p><b>Free block bitmap</b> — File names usually given as a <b>path</b> from root directory. Relative pathnames are relative to current working directory.</p> <p><b>Data and inode blocks</b> — Relative pathnames based off current working directory.</p> <p><b>Directory Operations</b> — <i>Open/Close, Search, Create/Delete, Link/Unlink, Change directory, List, Read/Write attributes, Mount</i>.</p> <p><b>LINK</b> — reference to directory/file in another part of FS.</p> <p><b>Hard link</b> — references address, <b>symbolic (soft) link</b> — references file name.</p> <p><b>MOUNT</b> operation combines multiple FSs into one namespace, allows references from single root dir. <b>Mount point</b> is the dir in native FS assigned to root of mounted FS. FSs manage mounted dirs with <b>mount tables</b>.</p> <p><b>EXT2FS</b> is a high-performance, robust FS with support for large files. File blocks are 4KB (4096/8192 bytes, with 5% of blocks reserved for root).</p> <p><b>ext2 inode</b> represents files/dirs. Provides fast access to small files, while supporting very large files.</p> <p><b>Block groups</b> are clusters of contiguous blocks — FS attempts to store related data in same block group to reduce seek time.</p> <p></p> <p><b>SECURITY</b> aims to prevent unauthorised access to the system, and permit authorised sharing of resources. Data <b>confidentiality, integrity &amp; system availability</b>. Security aspects — <b>system is as secure as weakest link</b>.</p> <p><b>PEOPLE SECURITY</b> — Large number of computer crime by insiders. Social engineering. May work around security measures for convenience. May have unrealistic security expectations (e.g. being unable to forge email addresses).</p> <p><b>HARDWARE SECURITY</b> — With <b>physical access</b> to hardware, one can gain control of memory/disk, listen to network traffic (incl. passwords), forge messages on network. Exploitable security flaws.</p> <p><b>SOFTWARE SECURITY</b> — Software bugs may allow attacks to compromise system. Can gain root privileges, corrupt memory, steal data, damage data integrity and deny access to the system. May exploit buffer/integer overflows and format string vulnerabilities.</p> <p><b>AUTHENTICATION</b> — verify identity of users (<b>principals</b>) based on personal characteristics, possessions, knowledge. Personal characteristics hard to forge. Can suffer from equipment cost, false positives/negatives.</p>